

---

## **Report for lab1, Kexing Zhou**

## Contents

<b>Environment Configuration</b>	<b>2</b>
Test Compiler Toolchain . . . . .	2
QEMU Emulator . . . . .	3
<b>PC bootstrap</b>	<b>3</b>
Exercise 1 . . . . .	3
Exercise 2 . . . . .	3
Exercise 3 . . . . .	4
Exercise 4 . . . . .	5
Exercise 5 . . . . .	5
Exercise 6 . . . . .	6
Exercise 7 . . . . .	7
Exercise 8 . . . . .	7
Question 1 . . . . .	7
Question 2 . . . . .	8
Question 3 . . . . .	8
Question 4 . . . . .	8
Question 5 . . . . .	9
Question 6 . . . . .	9
Exercise 9 . . . . .	9
Exercise 10 . . . . .	10
Exercise 11 . . . . .	10
Exercise 12 . . . . .	11
Challenge 1 . . . . .	12

[TOC]

## Environment Configuration

```

1 Hardware Environment:
2 Memory:          16GB
3 Processor:       Intel Core i7-8550U CPU @ 1.66GHz 8
4 GPU:            NVIDIA GeForce RTX 2070
5 OS Type:        64 bit
6 Disk:           924GB
7
8 Software Environment:
9 OS:              Arch Linux
10 Gcc:             Gcc 11.1.0
11 Make:           GNU Make 4.3
12 Gdb:            GNU gdb 11.1

```

## Test Compiler Toolchain

```

1 $ objdump -i # the 5th line say elf32-i386
2 $ gcc -m32 -print-libgcc-file-name
3 /usr/lib/gcc/x86_64-pc-linux-gnu/11.1.0/32/libgcc.a

```

## QEMU Emulator

```
1 $ sudo pacman -S riscv64-linux-gnu-binutils \
2   riscv64-linux-gnu-gcc riscv64-linux-gnu-gdb qemu-arch-extra
```

## PC bootstrap

### Exercise 1

(nothing to report)

### Exercise 2

System starts at memory address 0xffff0. There, computer will read a `ljmp` command and jump to bios program start at 0xfe05b.

```
1 [f000:fff0] 0xffff0: ljmp $0x3630,$0xf000e05b
2 (gdb) si
3 [f000:e05b] 0xfe05b: cmpw $0x8,%cs:(%esi)
```

Bios first setup ss, sp, dx:

```
1 [f000:e062] 0xfe062: jne 0xd241d0a4
2 [f000:e066] 0xfe066: xor %edx,%edx
3 [f000:e068] 0xfe068: mov %edx,%ss
4 [f000:e06a] 0xfe06a: mov $0x7000,%sp
5 [f000:e070] 0xfe070: mov $0x33c,%dx
6 [f000:e076] 0xfe076: jmp 0x5576cf2a
```

Then setup interrupt, Global Descriptor Table and Interrupts Descriptor Table then launch protected mode.

```
1 [f000:cf28] 0xfc28: cli
2 [f000:cf29] 0xfc29: cld
3 [f000:cf2a] 0xfc2a: mov %ax,%cx
4 [f000:cf2d] 0xfc2d: mov $0x8f,%ax
5 # setup non-maskable interrupt
6 [f000:cf33] 0xfc33: out %al,$0x70
7 [f000:cf35] 0xfc35: in $0x71,%al
8 # setup system devices
9 [f000:cf37] 0xfc37: in $0x92,%al
10 [f000:cf39] 0xfc39: or $0x2,%al
11 [f000:cf3b] 0xfc3b: out %al,$0x92
12 [f000:cf3d] 0xfc3d: mov %cx,%ax
13 # setup GDT and IDT
14 [f000:cf40] 0xfc40: lidt %cs:(%esi)
15 [f000:cf46] 0xfc46: lgdt %cs:(%esi)
16 # launch protected mode
17 [f000:cf4c] 0xfc4c: mov %cr0,%ecx
18 [f000:cf4f] 0xfc4f: and $0xffff,%cx
19 [f000:cf56] 0xfc56: or $0x1,%cx
20 [f000:cf5a] 0xfc5a: mov %ecx,%cr0
21 # goto 32bit code
22 [f000:cf5d] 0xfc5d: ljmpw $0xf,$0xcf65
```

Then bios setup ds, es, ss, fs, gs, and jump to bios code.

```
1 0xfc65: mov $0x10,%ecx
2 0xfc6a: mov %ecx,%ds
3 0xfc6c: mov %ecx,%es
4 0xfc6e: mov %ecx,%ss
5 0xfc70: mov %ecx,%fs
6 0xfc72: mov %ecx,%gs
7 0xfc74: jmp *%edx
```

The bios code is too complex to analysis, but it looks like it is well-compiled by some compiler:

```
1 (gdb) x/10i
2 0xf033c: push %ebx
```

```
3 0xf033d: sub $0x20,%esp
4 0xf0340: call 0xecd04
5 0xf0345: mov $0x40000000,%ebx
6 0xf034a: lea 0xc(%esp),%eax
7 0xf034e: push %eax
8 0xf034f: lea 0xc(%esp),%eax
9 0xf0353: push %eax
10 0xf0354: lea 0xc(%esp),%ecx
11 0xf0358: lea 0x8(%esp),%edx
12 0xf035c: mov %ebx,%eax
```

I think bios set up some system device in these process, and exit protected mode and jump to start point, 0x7c00.

I tried `watch $cr0` to find out when the bios exits protected mode, but it stuck gdb and not works well. The exit of protected mode is very misterious.

### Exercise 3

The disas code at 0x7c00 is:

```
1 (gdb) b *0x7c00
2 Breakpoint 1 at 0x7c00
3 (gdb) c
4 Continuing.
5 [ 0:7c00] => 0x7c00: cli
6
7 Breakpoint 1, 0x00007c00 in ?? ()
8 (gdb) x/10i
9 0x7c01: cld
10 0x7c02: xor %eax,%eax
11 0x7c04: mov %eax,%ds
12 0x7c06: mov %eax,%es
13 0x7c08: mov %eax,%ss
14 0x7c0a: in $0x64,%al
15 0x7c0c: test $0x2,%al
16 0x7c0e: jne 0x7c0a
17 0x7c10: mov $0xd1,%al
18 0x7c12: out %al,$0x64
19 (gdb)
```

The assembly code are equal in boot.S and boot.asm.

**Identify the exact assembly instructions that correspond to each of the statements in `readsect()`:**

At boot.asm, Line 150:

```
1 00007c78 <readsect>:
2
3 void
4 readsect(void *dst, uint32_t offset)
5 {
6     7c78: 55                push %ebp
7     7c79: 89 e5            mov %esp,%ebp
8     7c7b: 57              push %edi
9     7c7c: 50              push %eax
10    7c7d: 8b 4d 0c         mov 0xc(%ebp),%ecx
11    // wait for disk to be ready
12    waitdisk();
13    7c80: e8 e5 ff ff     call 7c6a <waitdisk>
14 }
```

**identify the begin and end of the for loop that reads the remaining sectors of the kernel from the disk.**

At boot.asm, Line 291.

```
1 for (; ph < eph; ph++)
2 7d56: 39 f3            cmp %esi,%ebx
3 7d58: 73 17            jae 7d71 <bootmain+0x58>
4     readseg(ph->p_pa, ph->p_memsz, ph->p_offset);
5 7d5a: 50              push %eax
6 for (; ph < eph; ph++)
7 7d5b: 83 c3 20         add $0x20,%ebx
8     readseg(ph->p_pa, ph->p_memsz, ph->p_offset);
9 7d5e: ff 73 e4         push -0x1c(%ebx)
```

```

10 7d61: ff 73 f4      push    -0xc(%ebx)
11 7d64: ff 73 ec      push    -0x14(%ebx)
12 7d67: e8 6e ff ff   call    7cda <readseg>
13 for (; ph < eph; ph++)
14 7d6c: 83 c4 10      add     $0x10,%esp
15 7d6f: eb e5        jmp     7d56 <bootmain+0x3d>
16 ((void (*)(void)) (ELFHDR->e_entry))();
17 7d71: ff 15 18 00 01 00 call    *0x10018

```

set a breakpoint there, and continue to that breakpoint.

```

1 (gdb) b *(0x7d71)
2 Breakpoint 2 at 0x7d71
3 (gdb) c
4 Continuing.
5 => 0x7d71:      call    *0x10018
6 (gdb) si
7 => 0x10000c:     movw    $0x1234,0x472
8 0x0010000c in ?? ()
9 (gdb) x/10i
10 0x100015:     mov     $0x111000,%eax
11 0x10001a:     mov     %eax,%cr3
12 0x10001d:     mov     %cr0,%eax
13 0x100020:     or      $0x80010001,%eax
14 0x100025:     mov     %eax,%cr0
15 0x100028:     mov     $0xf010002f,%eax
16 0x10002d:     jmp     *%eax
17 0x10002f:     mov     $0x0,%ebp
18 0x100034:     mov     $0xf010f000,%esp
19 0x100039:     call    0x1000a6

```

## Exercise 4

```

1 void
2 f(void)
3 {
4     int a[4];
5     int *b = malloc(16);
6     int *c;
7     int i;
8
9     // print address of a, b, c
10    printf("1: a = %p, b = %p, c = %p\n", a, b, c);
11
12    c = a; // make c point to a[1]
13    for (i = 0; i < 4; i++)
14        a[i] = 100 + i; // set the i-th elements of a 100 + i
15    c[0] = 200; // change c[0] to 200, a[0] also changes
16    // print value of array A's first to forth element
17    printf("2: a[0] = %d, a[1] = %d, a[2] = %d, a[3] = %d\n",
18          a[0], a[1], a[2], a[3]);
19    // value in a: 200 101 102 103
20
21    c[1] = 300; // change c[1] to 300, a[1] also changes
22    *(c + 2) = 301; // change *(c+2) = c[2] to 301, a[2] also changes
23    3[c] = 302; // 3[c] mean *(3 + c) = *(c + 3) = c[3], change c[3] to 302, a[3] also changes
24    printf("3: a[0] = %d, a[1] = %d, a[2] = %d, a[3] = %d\n", #
25          a[0], a[1], a[2], a[3]);
26    // value in a: 200 300 301 302
27
28    c = c + 1; // make c point to a[1]
29    *c = 400; // change c[0] to 400, a[1] also changes
30    printf("4: a[0] = %d, a[1] = %d, a[2] = %d, a[3] = %d\n",
31          a[0], a[1], a[2], a[3]);
32    // value in a: 200 400 301 302
33
34    c = (int *) ((char *) c + 1); // make c point to a byte after a[1]
35    *c = 500; // change the last 3 byte of a[1] and the first byte of a[2]
36    printf("5: a[0] = %d, a[1] = %d, a[2] = %d, a[3] = %d\n",
37          a[0], a[1], a[2], a[3]);
38    // value in a: 200 128144 256 302
39
40    b = (int *) a + 1; // make b point to a[1]
41    c = (int *) ((char *) a + 1); // make c point to a byte after a[0]
42    printf("6: a = %p, b = %p, c = %p\n", a, b, c);
43 }

```

## Exercise 5

The original bootloader:

```

1 => 0x7c00:     cli
2 0x7c01:     cld
3 0x7c02:     xor     %eax,%eax

```

```

4  0x7c04:    mov    %eax,%ds
5  0x7c06:    mov    %eax,%es
6  0x7c08:    mov    %eax,%ss
7  0x7c0a:    in     $0x64,%al
8  0x7c0c:    test   $0x2,%al
9  0x7c0e:    jne     0x7c0a
10 0x7c10:    mov    $0xd1,%al
11 0x7c12:    out    %al,$0x64
12 0x7c14:    in     $0x64,%al
13 0x7c16:    test   $0x2,%al
14 0x7c18:    jne     0x7c14
15 0x7c1a:    mov    $0xdf,%al
16 0x7c1c:    out    %al,$0x60
17 0x7c1e:    lgdtl   (%esi)
18 0x7c21:    fs    jl 0x7c33
19 0x7c24:    and     %al,%al
20 0x7c26:    or      $0x1,%ax
21 0x7c2a:    mov     %eax,%cr0
22 0x7c2d:    ljmp    $0xb866,$0x87c32 # (*)
23 0x7c34:    adc     %al,(%eax)
24 0x7c36:    mov     %eax,%ds
25 0x7c38:    mov     %eax,%es
26 0x7c3a:    mov     %eax,%fs
27 0x7c3c:    mov     %eax,%gs
28 0x7c3e:    mov     %eax,%ss
29 0x7c40:    mov     $0x7c00,%esp
30 0x7c45:    call    0x7d19

```

I change boot/Makefrag, Line 28:

```

1 - $(V)$(LD) $(LDFLAGS) -N -e start -Ttext 0x7c00 -o $@.out $^
2 + $(V)$(LD) $(LDFLAGS) -N -e start -Ttext 0x0000 -o $@.out $^

```

Then restart the bootloader.

```

1 => 0x7c00:    cli
2  0x7c01:    cld
3  0x7c02:    xor     %eax,%eax
4  0x7c04:    mov     %eax,%ds
5  0x7c06:    mov     %eax,%es
6  0x7c08:    mov     %eax,%ss
7  0x7c0a:    in      $0x64,%al
8  0x7c0c:    test    $0x2,%al
9  0x7c0e:    jne     0x7c0a
10 0x7c10:    mov     $0xd1,%al
11 0x7c12:    out     %al,$0x64
12 0x7c14:    in      $0x64,%al
13 0x7c16:    test    $0x2,%al
14 0x7c18:    jne     0x7c14
15 0x7c1a:    mov     $0xdf,%al
16 0x7c1c:    out     %al,$0x60
17 0x7c1e:    lgdtl   (%esi)
18 0x7c21:    add     %cl,%fs:(%edi)
19 0x7c24:    and     %al,%al
20 0x7c26:    or      $0x1,%ax
21 0x7c2a:    mov     %eax,%cr0
22 0x7c2d:    ljmp    $0xb866,$0x80032 # (*)
23 0x7c34:    adc     %al,(%eax)
24 0x7c36:    mov     %eax,%ds
25 0x7c38:    mov     %eax,%es
26 0x7c3a:    mov     %eax,%fs
27 0x7c3c:    mov     %eax,%gs
28 0x7c3e:    mov     %eax,%ss
29 0x7c40:    mov     $0x0,%esp
30 0x7c45:    call    0x7d19

```

As seen, command are all the same, and short jump target address are the same, also. This is because command type is independent from the code address, and short jump is PC-relative addressing. But long jump target address is changed to a wrong place. So, the modified bootloader fails to load the kernel.

## Exercise 6

```

1 (gdb) b *0x7c00 # break at bootloader
2 Breakpoint 1 at 0x7c00
3 (gdb) c
4 Continuing.
5 [ 0x7c00] => 0x7c00: cli
6
7 Breakpoint 1, 0x00007c00 in ?? ()
8 (gdb) x/8w 0x00100000
9 0x100000: 0x00000000 0x00000000 0x00000000 0x00000000
10 0x100010: 0x00000000 0x00000000 0x00000000 0x00000000
11 (gdb) b *0x7d71 # break at enter kernel

```

```
12 Breakpoint 2 at 0x7d71
13 (gdb) c
14 Continuing.
15 The target architecture is set to "i386".
16 => 0x7d71:      call    *0x10018
17
18 Breakpoint 2, 0x00007d71 in ?? ()
19 (gdb) x/8w 0x00100000
20 0x100000:      0x1badb002      0x00000000      0xe4524ffe      0x7205c766
21 0x100010:      0x34000004      0x1000b812      0x220f0011      0xc0200fd8
```

The address 0x00100000 is the kernel start point. When bios enters the boot loader, the kernel is not loaded, so they are all set to zero. The value changes when the boot loader loads kernel from disk.

## Exercise 7

The code `movl %eax, %cr0` is located at Line 37 in `obj/kern/kernel.asm`. Then, I break at the load address 0x100025.

```
1 (gdb) b *0x100025
2 (gdb) c
3 (gdb) x/10w 0x00100000
4 0x100000:      0x1badb002      0x00000000      0xe4524ffe      0x7205c766
5 0x100010:      0x34000004      0x1000b812      0x220f0011      0xc0200fd8
6 0x100020:      0x0100010d      0xc0220f80
7 (gdb) x/10w 0xf0100000
8 0xf0100000 <_start-268435468>: Cannot access memory at address 0xf0100000
9 (gdb) si
10 (gdb) x/10w 0x00100000
11 0x100000:      0x1badb002      0x00000000      0xe4524ffe      0x7205c766
12 0x100010:      0x34000004      0x1000b812      0x220f0011      0xc0200fd8
13 0x100020:      0x0100010d      0xc0220f80
14 (gdb) x/10w 0xf0100000
15 0xf0100000 <_start-268435468>: 0x1badb002      0x00000000      0xe4524ffe      0x7205c766
16 0xf0100010 <entry+4>:      0x34000004      0x1000b812      0x220f0011      0xc0200fd8
17 0xf0100020 <entry+20>: 0x0100010d      0xc0220f80
```

As seen, the data at 0x00100000 doesn't change, because we map virtual memory [0, 4Mb) to physical memory [0, 4Mb). But the data at 0xf0100000 changed, and is the same as that in 0x00100000, because we map [0xf0100000, 0xffffffff) to [0, 4Mb). After virtual memory translation, they are the same physical address.

## Comment out the `movl %eax, %cr0` in `kern/entry.S`, trace into it

```
1 (gdb) b *0x100020 # the previous code
2 (gdb) c
3 (gdb) x/10w 0x00100000
4 0x100000:      0x1badb002      0x00000000      0xe4524ffe      0x7205c766
5 0x100010:      0x34000004      0x1000b812      0x220f0011      0xc0200fd8
6 0x100020:      0x0100010d      0x002cb880
7 (gdb) x/10w 0xf0100000
8 0xf0100000 <_start-268435468>: Cannot access memory at address 0xf0100000
9 (gdb) si
10 (gdb) x/10w 0x00100000
11 0x100000:      0x1badb002      0x00000000      0xe4524ffe      0x7205c766
12 0x100010:      0x34000004      0x1000b812      0x220f0011      0xc0200fd8
13 0x100020:      0x0100010d      0x002cb880
14 (gdb) x/10w 0xf0100000
15 0xf0100000 <_start-268435468>: Cannot access memory at address 0xf0100000
```

As seen, the result doesn't change, because we don't enable virtual memory at all.

## Exercise 8

### Question 1

**Explain the interface between `printf.c` and `console.c`. Specifically, what function does `console.c` export? How is this function used by `printf.c`?**

console.c exports cputchar. This function hides the details of the interaction with the hardware and only requires cprintf.c to give it a character to output it.

## Question 2

### Explain the following from console.c

This part of the code checks if the character buffer on the screen is full, and if it is, deletes the first line, moves the other lines forward, and set the last line to blank (0x0700 means console style is black on white). To achieve the screen scrolling effect.

## Question 3

### In the call to cprintf(), to what does fmt point? To what does ap point?

fmt points "x %d, y %x, z %d\n", ap point to x in the stack frame.

**List (in order of execution) each call to cons\_putc, va\_arg, and vprintf. For cons\_putc, list its argument as well. For va\_arg, list what ap points to before and after the call. For vprintf list the values of its two arguments.**

```
1  vprintf("x %d, y %x, z %d\n", ap) ap point to x
2  cons_putc('x')
3  cons_putc(' ')
4  va_arg, before ap->x, after ap->y
5  cons_putc('1')
6  cons_putc(',')
7  cons_putc(' ')
8  cons_putc('y')
9  cons_putc(' ')
10 va_arg, before ap->y, after ap->z
11 cons_putc('3')
12 cons_putc(',')
13 cons_putc(' ')
14 cons_putc('z')
15 cons_putc(' ')
16 va_arg, before ap->z, after ap->(unknow address after z)
17 cons_putc('4')
18 cons_putc('\n')
```

## Question 4

### What is the output?

He110 World

**Explain how this output is arrived at in the step-by-step manner of the previous exercise.**

```
1  vprintf("H%x Wo%s", ap), ap point to 57616
2  cons_putc('H')
3  va_arg, before ap->57616, after ap->the value &i in stack
4  cons_putc('e')
5  cons_putc('1')
6  cons_putc('1')
7  cons_putc('0')
8  cons_putc(' ')
9  cons_putc('W')
10 cons_putc('o')
11 va_arg, before ap->the value &i in stack, after ap->(unknow address after the value &i in the stack)
12 cons_putc('r')
13 cons_putc('l')
14 cons_putc('d')
```

In little-endian, variable i is in memory as



```
1 0x72 0x6c 0x64 0x00
```

When we use %s to print the string at &i, we decode the above value into ASCII characters, so the string is:

```
1 rld\0
```

The symbol \0 is a identifier for string ending.

**If the x86 were instead big-endian what would you set i to in order to yield the same output?**  
0x726c6400

**Would you need to change 57616 to a different value?** No

### Question 5

**In the following code, what is going to be printed after 'y='? (note: the answer is not a specific value.)**

y will print the integer value in the stack just above variable x.

**Why does this happen?**

```
1  vcprintf
2  cons_putc('x')
3  cons_putc('=')
4  va_arg, before ap->x, after ap->(the stack value after x)
5  cons_putc('3')
6  cons_putc(' ')
7  cons_putc('y')
8  cons_putc('=')
9  va_arg, before ap->(the stack value after x), after ap->(the atack value 4byte above x)
10 .....(it will print a strange value)
```

### Question 6

**Let's say that GCC changed its calling convention so that it pushed arguments on the stack in declaration order, so that the last argument is pushed last. How would you have to change cprintf or its interface so that it would still be possible to pass it a variable number of arguments?**

Change the position of parameter `fmt` to the last position. Then, `fmt` always just above the place of return address in the stack frame. We find the position of `fmt` first, then parse the parameter list, then use `va_arg` to access the variable reversely.

### Exercise 9

**Determine where the kernel initializes its stack**

At kern/entry.S, Line 77:

```
1  movl    $(bootstacktop),%esp
```

**and exactly where in memory its stack is located**

```
1  $ readelf obj/kern/kernel -a | grep stack
2  106: f0110000      0 NOTYPE  GLOBAL DEFAULT   6 bootstacktop
3  109: f0108000      0 NOTYPE  GLOBAL DEFAULT   6 bootstack
```

As seen, the stack is located at 0xf0108000(virtual memory), 0x00108000(physical memory).

### How does the kernel reserve space for its stack?

At kern/entry.S, Line 92:

```
1  bootstack:
2      .space      KSTKSIZE
3      .globl      bootstacktop
4  bootstacktop:
```

The command `.space` reserves `KSTKSIZE` space in the program.

### And at which "end" of this reserved area is the stack pointer initialized to point to?

It points to the end with highest address, because stack in C grows downward.

## Exercise 10

To become familiar with the C calling conventions on the x86, find the address of the `test_backtrace` function in `obj/kern/kernel.asm`, set a breakpoint there, and examine what happens each time it gets called after the kernel starts. How many 32-bit words does each recursive nesting level of `test_backtrace` push on the stack, and what are those words?

```
1  (gdb) b test_backtrace
2  Breakpoint 1 at 0xf0100040: file kern/init.c, line 13.
3  (gdb) c
4  (gdb) c 5
5  Will ignore next 4 crossings of breakpoint 1. Continuing.
6  => 0xf0100040 <test_backtrace>: push %ebp
7  Breakpoint 1, test_backtrace (x=0) at kern/init.c:13
8  (gdb) x/36wd $sp
9  (gdb) x/36wx $sp
10 ; RET addr of test_backtrace(0) parameters ??? ???
11 0xf010ef3c: 0xf0100076 0x00000000 0x00000001 0xf010ef78
12 ; ??? ??? ??? ebp
13 0xf010ef4c: 0xf010004a 0xf0110308 0x00000002 0xf010ef78
14 ; RET addr of test_backtrace(1) parameters ??? ???
15 0xf010ef5c: 0xf0100076 0x00000001 0x00000002 0xf010ef98
16 ; ??? ??? ??? ebp
17 0xf010ef6c: 0xf010004a 0xf0110308 0x00000003 0xf010ef98
18 ; RET addr of test_backtrace(2) parameters ??? ???
19 0xf010ef7c: 0xf0100076 0x00000002 0x00000003 0xf010efb8
20 ; ??? ??? ??? ebp
21 0xf010ef8c: 0xf010004a 0xf0110308 0x00000004 0xf010efb8
22 ; RET addr of test_backtrace(3) parameters ??? ???
23 0xf010ef9c: 0xf0100076 0x00000003 0x00000004 0x00000000
24 ; ??? ??? ??? ebp
25 0xf010efac: 0xf010004a 0xf0110308 0x00000005 0xf010efd8
26 ; RET addr of test_backtrace(4) parameters ??? ???
27 0xf010efbc: 0xf0100076 0x00000004 0x00000005 0x00000000
28 ; ??? ??? ??? ebp
29 0xf010efcc: 0xf010004a 0xf0110308 0x000100d4 0xf010eff8
30 ; RET addr of test_backtrace(5) parameters ??? ???
31 0xf010efdc: 0xf01000f4 0x00000005 0x00001aac 0x00000660
32 0xf010efec: 0x00000000 0x00000000 0x000100d4 0x00000000
33 0xf010effc: 0xf010003e 0x00000003 0x00001003 0x00002003
34 (gdb) bt
35 #0 test_backtrace (x=1) at kern/init.c:13
36 #1 0xf0100076 in test_backtrace (x=2) at kern/init.c:16
37 #2 0xf0100076 in test_backtrace (x=3) at kern/init.c:16
38 #3 0xf0100076 in test_backtrace (x=4) at kern/init.c:16
39 #4 0xf0100076 in test_backtrace (x=5) at kern/init.c:16
40 #5 0xf01000f4 in i386_init () at kern/init.c:39
41 #6 0xf010003e in relocated () at kern/entry.S:80
```

In the stack dump, we can see, at each recursive nesting level, `test_backtrace` push 8 32-bit words into the stack. They are return address, parameter `x`, and other temporary data.

## Exercise 11

At monitor.c Line 57:

```
1 int
2 mon_backtrace(int argc, char **argv, struct Trapframe *tf)
3 {
4     printf("Stack backtrace:\n");
5     uint32_t *ebp = (uint32_t*)read_ebp();
6     do {
7         printf("ebp %08x eip %08x args %08x %08x %08x %08x\n",
8             ebp, *(ebp + 1), *(ebp + 2), *(ebp + 3), *(ebp + 4), *(ebp + 5));
9         ebp = (uint32_t*)*ebp;
10    } while(ebp);
11    printf("ebp %08x eip %08x args %08x %08x %08x %08x\n",
12        ebp, *(ebp + 1), *(ebp + 2), *(ebp + 3), *(ebp + 4), *(ebp + 5));
13    return 0;
14 }
```

The loop ends when prevebp equals to zero, because in entry.S, Line 69, the ebp's initial value is zero.

```
1 relocated:
2
3     # Clear the frame pointer register (EBP)
4     # so that once we get into debugging C code,
5     # stack backtraces will be terminated properly.
6     movl    $0x0,%ebp    # nuke frame pointer
```

**The return instruction pointer typically points to the instruction after the call instruction (why?)**

The command `return` is equals to `popl eip`, so we need to save the next command, which runs after function calls, into stack.

**(Why can't the backtrace code detect how many arguments there actually are? How could this limitation be fixed?)**

The memory space for arguments and local variable are in the same stack frame. Without debug info, we cannot figure out which are arguments and which are local variables. So we can't detect how many arguments it is. We can search debug message to determine the argument count.

## Exercise 12

In `debuginfo_eip`, where do `_STAB` come from?\*

Compiler transforms code into assembly code. When stab table is enabled, the assembly code will contain `.stab` command to tell linker the symbol and its info.

**Complete the implementation of `debuginfo_eip` by inserting the call to `stab_binsearch` to find the line number for an address.**

monitor.c, Line 59

```
1 int
2 mon_backtrace(int argc, char **argv, struct Trapframe *tf)
3 {
4     printf("Stack backtrace:\n");
5     struct Eipdebuginfo info;
6     volatile uint32_t *ebp = (uint32_t*)read_ebp(), eip;
7     do {
8         eip = *(ebp + 1);
9         debuginfo_eip(eip - 4, &info);
10        printf("  ebp %08x eip %08x args %08x %08x %08x %08x\n",
11            ebp, eip, *(ebp + 2), *(ebp + 3), *(ebp + 4), *(ebp + 5));
12        printf("    %s:%d: ", info.eip_file, info.eip_line);
13        for(int i = 0; i < info.eip_fn_namelen; i++) {
14            putchar(info.eip_fn_name[i]);
15        }
16        printf("%d\n", eip - info.eip_fn_addr);
17        if(ebp == NULL) break;
18        ebp = (uint32_t*)*ebp;
19    } while(ebp);
20    return 0;
21 }
```

kdebug.c, line 182

```
1 stab_binsearch(stabs, &lline, &rline, N_SLINE, addr);
2 info->eip_line = stabs[lline].n_desc;
```

## Challenge 1

The reference in the lab page is incomplete, I found a complete reference for color in wiki.

I implemented some ansi feature, include:

- set\_cursor\_pos
- cursor\_up, down, forward, backward
- save\_cursor\_position, restore\_cursor\_position
- erase\_display, erase\_line
- set\_graphical\_mode (only set color is implemented)

You can test my code, runs:


```
1 $ make qemu
2 K> colortest
```



The result is:

The clear command is also implemented. The result is:

```
1 $ make qemu
2 K> clear
```



To the terminal display color correctly, the code of color escaping should be put into cga module. So it needs a context. The code is a bit long.

```
1 static unsigned addr_6845;
2 static uint16_t *crt_buf;
3 static uint16_t crt_pos;
4
5 struct ANSIText {
6     int mode;
7     bool last_esc;
8     bool is_escaping;
9     int args[8];
10    int acnt;
11    uint16_t saved_pos;
12 } crt_ctx;
13
14 void cga_set_cursor_pos(int PL, int Pc) {
15     uint16_t target = PL * CRT_COLS + Pc;
16     if(target < CRT_SIZE) crt_pos = target;
17 }
18
19 void cga_cursor_up(int Pn) {
20     uint16_t target = crt_pos - Pn * CRT_COLS;
21     if(target < CRT_SIZE) crt_pos = target;
22 }
23
24 void cga_cursor_down(int Pn) {
25     uint16_t target = crt_pos + Pn * CRT_COLS;
26     if(target < CRT_SIZE) crt_pos = target;
27 }
28
29 void cga_cursor_backward(int Pn) {
30     uint16_t target = crt_pos - Pn;
31     if(target < CRT_SIZE) crt_pos = target;
32 }
33
34 void cga_cursor_forward(int Pn) {
35     uint16_t target = crt_pos + Pn;
36     if(target < CRT_SIZE) crt_pos = target;
37 }
38
39 void cga_save_cursor_position() {
40     crt_ctx.saved_pos = crt_pos;
41 }
42
43 void cga_restore_cursor_position() {
44     uint16_t target = crt_ctx.saved_pos;
45     if(target < CRT_SIZE) crt_pos = target;
46 }
47
```

```

48 void cga_erase_display() {
49     for(uint16_t i = 0; i < CRT_SIZE; i++) {
50         crt_buf[i] = crt_ctx.mode | ' ';
51     }
52     crt_pos = 0;
53 }
54
55 void cga_erase_line() {
56     uint16_t ed = crt_pos - crt_pos % CRT_ROWS + CRT_ROWS;
57     for(uint16_t i = crt_pos; i < ed; i++) {
58         crt_buf[i] = crt_ctx.mode | ' ';
59     }
60 }
61
62 void cga_set_graphical_mode(int modecode) {
63     if(0 <= modecode && modecode < 8) {
64         switch(modecode) {
65             case 0: crt_ctx.mode = 0x0700; break;
66         }
67     }
68     else if(30 <= modecode && modecode < 38) {
69         modecode -= 30;
70         crt_ctx.mode &= 0xf000;
71         crt_ctx.mode |= modecode << 8;
72     }
73     else if(90 <= modecode && modecode < 98) {
74         modecode = modecode - 90 + 8;
75         crt_ctx.mode &= 0xf000;
76         crt_ctx.mode |= modecode << 8;
77     }
78     else if(40 <= modecode && modecode < 48) {
79         modecode -= 40;
80         crt_ctx.mode &= 0x0f00;
81         crt_ctx.mode |= modecode << 12;
82     }
83     else if(100 <= modecode && modecode < 108) {
84         modecode = modecode - 100 + 8;
85         crt_ctx.mode &= 0x0f00;
86         crt_ctx.mode |= modecode << 12;
87     }
88 }
89
90 void cga_set_mode(int modecode) {
91     panic("cga_set_mode not implemented");
92 }
93
94 void cga_reset_mode() {
95     panic("cga_reset_mode not implemented");
96 }
97
98 static void
99 cga_init(void)
100 {
101     volatile uint16_t *cp;
102     uint16_t was;
103     unsigned pos;
104
105     crt_ctx.mode = 0x0700;
106     crt_ctx.is_escaping = false;
107     crt_ctx.last_esc = false;
108
109     cp = (uint16_t*) (KERNBASE + CGA_BUF);
110     was = *cp;
111     *cp = (uint16_t) 0xA55A;
112     if (*cp != 0xA55A) {
113         cp = (uint16_t*) (KERNBASE + MONO_BUF);
114         addr_6845 = MONO_BASE;
115     } else {
116         *cp = was;
117         addr_6845 = CGA_BASE;
118     }
119
120     /* Extract cursor location */
121     outb(addr_6845, 14);
122     pos = inb(addr_6845 + 1) << 8;
123     outb(addr_6845, 15);
124     pos |= inb(addr_6845 + 1);
125
126     crt_buf = (uint16_t*) cp;
127     crt_pos = pos;
128 }
129
130 static void cga_putc_raw(int c) {
131     switch (c & 0xff) {
132         case '\\b':
133             if (crt_pos > 0) {
134                 crt_pos--;
135                 crt_buf[crt_pos] = (c & ~0xff) | ' ';
136             }
137             break;
138         case '\\n':
139             crt_pos += CRT_COLS;
140             /* fallthru */
141         case '\\r':
142             crt_pos -= (crt_pos % CRT_COLS);
143             break;
144         case '\\t':
145             cons_putc(' ');
146             cons_putc(' ');

```

```

147     cons_putc(' ');
148     cons_putc(' ');
149     cons_putc(' ');
150     break;
151 default:
152     crt_buf[crt_pos++] = c;    /* write the character */
153     break;
154 }
155
156 // What is the purpose of this?
157 if (crt_pos >= CRT_SIZE) {
158     int i;
159
160     memmove(crt_buf, crt_buf + CRT_COLS, (CRT_SIZE - CRT_COLS) * sizeof(uint16_t));
161     for (i = CRT_SIZE - CRT_COLS; i < CRT_SIZE; i++)
162         crt_buf[i] = 0x0700 | ' ';
163     crt_pos -= CRT_COLS;
164 }
165 }
166
167 static void
168 cga_putc(int c) {
169     c &= 0xff;
170     if (c == 0x1b) {
171         crt_ctx.last_esc = true;
172     }
173     else if (crt_ctx.last_esc) {
174         if (c == '[') {
175             crt_ctx.last_esc = false;
176             crt_ctx.is_escaping = true;
177             crt_ctx.acnt = 0;
178             crt_ctx.args[0] = 0;
179         }
180         else {
181             crt_ctx.last_esc = false;
182             cga_putc_raw('\x1b');
183             cga_putc_raw('[');
184         }
185     }
186     else if (crt_ctx.is_escaping) {
187         switch (c) {
188             case '0': case '1': case '2': case '3': case '4':
189             case '5': case '6': case '7': case '8': case '9':
190                 crt_ctx.args[crt_ctx.acnt] = crt_ctx.args[crt_ctx.acnt] * 10 + c - '0';
191                 break;
192
193             case ';':
194                 if (crt_ctx.acnt < 7)
195                     crt_ctx.args[++crt_ctx.acnt] = 0;
196                 break;
197
198             case 'H':
199             case 'f': cga_set_cursor_pos(crt_ctx.args[0], crt_ctx.args[1]); goto crt_ansi_cmd_finish;
200             case 'A': cga_cursor_up(crt_ctx.args[0]); goto crt_ansi_cmd_finish;
201             case 'B': cga_cursor_down(crt_ctx.args[0]); goto crt_ansi_cmd_finish;
202             case 'C': cga_cursor_forward(crt_ctx.args[0]); goto crt_ansi_cmd_finish;
203             case 'D': cga_cursor_backward(crt_ctx.args[0]); goto crt_ansi_cmd_finish;
204             case 's': cga_save_cursor_position(); goto crt_ansi_cmd_finish;
205             case 'u': cga_restore_cursor_position(); goto crt_ansi_cmd_finish;
206             case 'J': cga_erase_display(); goto crt_ansi_cmd_finish;
207             case 'K': cga_erase_line(); goto crt_ansi_cmd_finish;
208             case 'm': for (int i = 0; i <= crt_ctx.acnt; i++) cga_set_graphical_mode(crt_ctx.args[i]); goto
209                 crt_ansi_cmd_finish;
210             case '=': goto crt_ansi_cmd_finish;
211             case 'h': cga_set_mode(crt_ctx.args[0]); goto crt_ansi_cmd_finish;
212             case 'l': cga_reset_mode(); goto crt_ansi_cmd_finish;
213
214             default:
215                 crt_ansi_cmd_finish:
216                 crt_ctx.is_escaping = false;
217         }
218     }
219     else {
220         crt_ctx.last_esc = false;
221         cga_putc_raw(c | crt_ctx.mode);
222     }
223     /* move that little blinky thing */
224     outb(addr_6845, 14);
225     outb(addr_6845 + 1, crt_pos >> 8);
226     outb(addr_6845, 15);
227     outb(addr_6845 + 1, crt_pos);

```