

---

**Report for lab4, Kexing Zhou,  
1900013008**

## Contents

<b>Environment Configuration</b>	<b>2</b>
Test Compiler Toolchain . . . . .	2
QEMU Emulator . . . . .	3
<b>Preemptive Multitasking</b>	<b>3</b>
Exercise 1 . . . . .	3
Exercise 2 . . . . .	3
Question 1 . . . . .	3
Exercise 3 . . . . .	3
Exercise 4 . . . . .	3
Exercise 5 . . . . .	4
Question 2 . . . . .	5
Exercise 6 . . . . .	5
Question 3 . . . . .	6
Challenge 3 . . . . .	6
Exercise 7 . . . . .	7
Exercise 8 . . . . .	8
Exercise 9 . . . . .	9
Exercise 10 . . . . .	9
Exercise 11 . . . . .	10
Exercise 12 . . . . .	10
Exercise 13 . . . . .	11
Exercise 14 . . . . .	11
Exercise 15 . . . . .	11
The end of this lab . . . . .	12

## Environment Configuration

1	Hardware Environment:
2	Memory: 16GB
3	Processor: Intel Core i7-8550U CPU @ 1.66GHz 8
4	GPU: NVIDIA GeForce RTX 2070
5	OS Type: 64 bit
6	Disk: 924GB
7	
8	Software Environment:
9	OS: Arch Linux
10	Gcc: Gcc 11.1.0
11	Make: GNU Make 4.3
12	Gdb: GNU gdb 11.1

## Test Compiler Toolchain

1	\$ objdump -i # the 5th line say elf32-i386
2	\$ gcc -m32 -print-libgcc-file-name
3	/usr/lib/gcc/x86_64-pc-linux-gnu/11.1.0/32/libgcc.a

## QEMU Emulator

```
1 $ sudo pacman -S riscv64-linux-gnu-binutils \
2   riscv64-linux-gnu-gcc riscv64-linux-gnu-gdb qemu-arch-extra
```

## Preemptive Multitasking

### Exercise 1

```
1 void *
2 mmio_map_region(physaddr_t pa, size_t size)
3 {
4     size = ROUNDUP(size, PGSIZE);
5     boot_map_region(kern_pgdir, base, size, pa, PTE_P | PTE_PCD | PTE_PWT | PTE_W);
6     uintptr_t ret = base;
7     base += size;
8     return (void*)ret;
9 }
```

### Exercise 2

```
1 for(size_t i = 1; i < npages_basemem; i++) {
2     if(i == PNUM(MPENTRY_PADDR)) continue; // This code ignores MPENTRY
3     pages[i].pp_ref = 0;
4     pages[i].pp_link = page_free_list;
5     page_free_list = &pages[i];
6 }
```

### Question 1

#### What is the purpose of macro MPBOOTPHYS?

The code which initialize MP is not linked at program position MPENTRY\_PADDR, but is loaded there. So the address of each variable need to be translated from the place it is linked to the MPENTRY\_PADDR. MPBOOTPHYS just did the translation statically.

#### Why is it necessary in kern/mpentry.S but not in boot/boot.S?

The `boot/boot.S` has already placed in 0x7000. It doesn't need the translation.

#### What could go wrong if it were omitted in kern/mpentry.S?

The address of each variable is wrong, then the MP fails to start up.

### Exercise 3

```
1 for(size_t i = 0; i < NCPU; i++) {
2     uintptr_t kstacktop_i = KSTACKTOP - i * (KSTACKSIZE + KSTKGAP);
3     boot_map_region(kern_pgdir, kstacktop_i - KSTACKSIZE, KSTACKSIZE, PADDR(percpu_kstacks[i]), PTE_P|PTE_W);
4 }
```

### Exercise 4

```
1 void
2 trap_init_percpu(void) {
3     int cpuid = cpunum();
4     struct Taskstate * ts = &(thiscpu->cpu_ts);
5
6     ts->ts_esp0 = KSTACKTOP - cpuid * (KSTKSIZE + KSTKGAP);
7     ts->ts_ss0 = GD_KD;
8     ts->ts_iomb = sizeof(struct Taskstate);
9
10    // Initialize the TSS slot of the gdt.
11    size_t gdt_idx = (GD_TSS0 >> 3) + cpuid;
12    gdt[gdt_idx] = SEG16(STS_T32A, (uint32_t) ts, sizeof(struct Taskstate) - 1, 0);
13    gdt[gdt_idx].sd_s = 0;
14
15    ltr(gdt_idx << 3);
16
17    // Load the IDT
18    lidt(&idt_pd);
19 }
```

## Exercise 5

The `lock_kernel` is called at:

```
1 void
2 i386_init(void) {
3     // .....
4     pic_init();
5     lock_kernel();
6     boot_aps();
7     // .....
8 }
```

```
1 void
2 mp_main(void) {
3     // .....
4     xchg(&thiscpu->cpu_status, CPU_STARTED);
5     lock_kernel();
6     sched_yield();
7     // .....
8 }
```

```
1 void
2 trap(struct Trapframe *tf)
3 {
4     // .....
5     // Re-acquire the big kernel lock if we were halted in
6     // sched_yield()
7     if (xchg(&thiscpu->cpu_status, CPU_STARTED) == CPU_HALTED)
8         lock_kernel();
9     assert(!(read_eflags() & FL_IF));
10
11     if ((tf->tf_cs & 3) == 3) {
12         // Trapped from user mode.
13         // Acquire the big kernel lock before doing any
14         // serious kernel work.
15         // LAB 4: Your code here.
16         assert(curenv);
17         lock_kernel();
18     }
19     // .....
20 }
```

The `unlock_kernel` is called at:

```
1 void
2 sched_halt(void) {
3     // .....
4
5     // Release the big kernel lock as if we were "leaving" the kernel
6     unlock_kernel();
7
8     // Reset stack pointer, enable interrupts and then halt.
9     asm volatile (
10         "movl $0, %%ebp\n"
11         "movl %0, %%esp\n"
12         "pushl $0\n" : : "r" (0));
13 }
```

```
1 void
2 env_run(struct Env *e) {
3     // LAB 3: Your code here.
4     assert(e->env_tf.tf_eflags & FL_IF);
5     unlock_kernel();
6     env_pop_tf(&(curenv->env_tf));
7 }
```

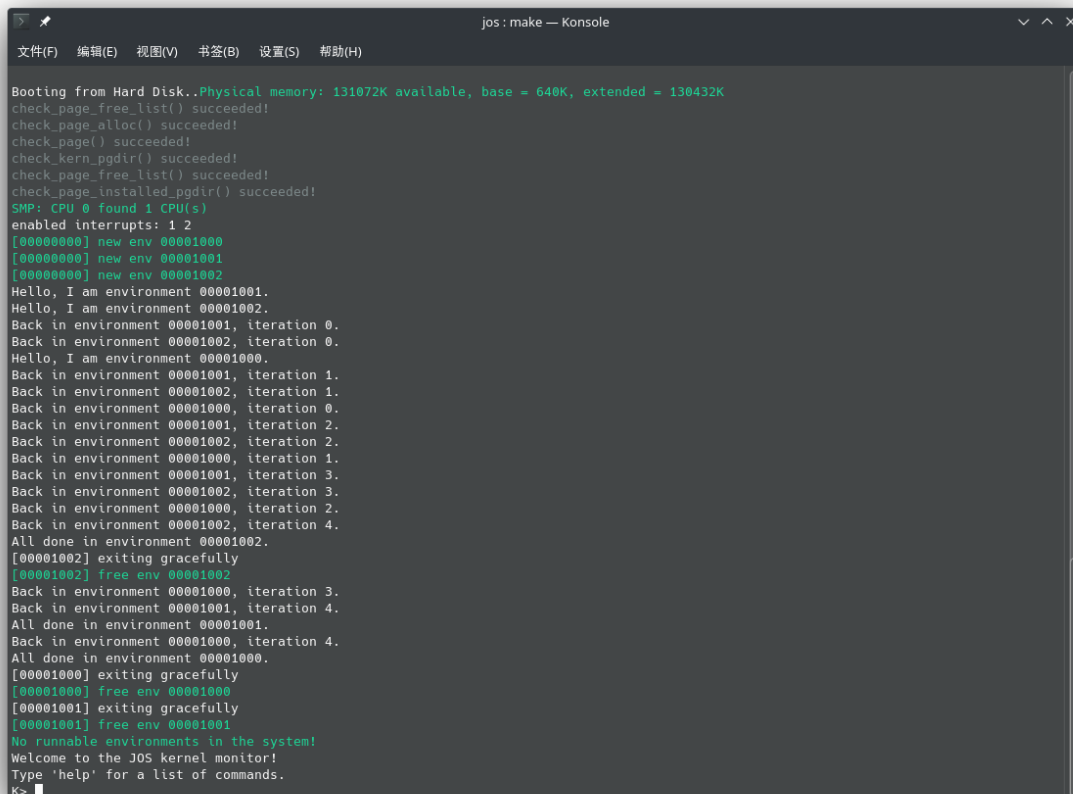
## Question 2

**It seems that using the big kernel lock guarantees that only one CPU can run the kernel code at a time. Why do we still need separate kernel stacks for each CPU? Describe a scenario in which using a shared kernel stack will go wrong, even with the protection of the big kernel lock.**

Two CPU receive exception at the same time. They both need to push execution context into memory, then acquires the big kernel lock. They share the stack so the context crashes. And the kernel fail to get execution context.

## Exercise 6

```
1 void
2 sched_yield(void) {
3     size_t env_id = curenv ? ENVX(curenv->env_id) : 0;
4     size_t i = env_id;
5     do {
6         if(envs[i].env_status == ENV_RUNNABLE) {
7             env_run(&envs[i]);
8         }
9         if(++i == NENV) i = 0;
10    } while(i != env_id);
11    if(curenv && (curenv->env_status == ENV_RUNNABLE || curenv->env_status == ENV_RUNNING)) {
12        env_run(curenv);
13    }
14    sched_halt();
15 }
```



```
jos: make — Konsole
文件(F) 编辑(E) 视图(V) 书签(B) 设置(S) 帮助(H)

Booting from Hard Disk..Physical memory: 131072K available, base = 640K, extended = 130432K
check_page_free_list() succeeded!
check_page_alloc() succeeded!
check_page() succeeded!
check_kern_pgdir() succeeded!
check_page_free_list() succeeded!
check_page_installed_pgdir() succeeded!
SMP: CPU 0 found 1 CPU(s)
enabled interrupts: 1 2
[00000000] new env 00001000
[00000000] new env 00001001
[00000000] new env 00001002
Hello, I am environment 00001001.
Hello, I am environment 00001002.
Back in environment 00001001, iteration 0.
Back in environment 00001002, iteration 0.
Hello, I am environment 00001000.
Back in environment 00001001, iteration 1.
Back in environment 00001002, iteration 1.
Back in environment 00001000, iteration 0.
Back in environment 00001001, iteration 2.
Back in environment 00001002, iteration 2.
Back in environment 00001000, iteration 1.
Back in environment 00001001, iteration 3.
Back in environment 00001002, iteration 3.
Back in environment 00001000, iteration 2.
Back in environment 00001002, iteration 4.
All done in environment 00001002.
[00001002] exiting gracefully
[00001002] free env 00001002
Back in environment 00001000, iteration 3.
Back in environment 00001001, iteration 4.
All done in environment 00001001.
Back in environment 00001000, iteration 4.
All done in environment 00001000.
[00001000] exiting gracefully
[00001000] free env 00001000
[00001001] exiting gracefully
[00001001] free env 00001001
No runnable environments in the system!
Welcome to the JOS kernel monitor!
Type 'help' for a list of commands.
K> 
```

### Question 3

#### Why can the pointer `e` be dereferenced both before and after the addressing switch?

Pointer `$e$` is stored in kernel stack, and where it points is in kernel memory. Since all environment share the same kernel memory mapping, the pointer `e` is dereferenced after addressing switch.

#### It must ensure the old environment's registers are saved so they can be restored properly later. Why?

The register contains context information such as stack pointer, program counter, and the current used data, which is critical for program running.

#### Where does this happen?

It happens in `kern/trap.c`, in function `trap`

```
1  if ((tf->tf_cs & 3) == 3) {
2      //.....
3      // Copy trap frame (which is currently on the stack)
4      // into 'curenv->env_tf', so that running the environment
5      // will restart at the trap point.
6      curenv->env_tf = *tf;
7      // The trapframe on the stack should be ignored from here on.
8      tf = &curenv->env_tf;
9  }
```

### Challenge 3

First, the FPU should be enabled:

In `entry.S` and `mpentry.S`:

```
1  # Turn on paging.
2  movl  %cr0, %eax
3  # CR0_MP enable x87 co-processor
4  orl  $(CR0_PE|CR0_PG|CR0_WP|CR0_MP), %eax
5  movl  %eax, %cr0
6
7  # CR4_OSFCSR enable FXSAVE and FXRSTOR instructions
8  # CR4_OSXMMEXCPT enable Unmasked SIMD Floating-Point Exceptions
9  movl  %cr4, %eax
10 orl  $(CR4_OSFCSR|CR4_OSXMMEXCPT), %eax
11 movl  %eax, %cr4
```

The command `fxsave` and `fxrstor` need 512byte aligned space to save the registers. So in `env.c`:

```
1  typedef uint8_t Fxbuf[512];
2  Fxbuf * fxbufs = NULL;
```

And `pmap.c`

```
1  fxbufs = boot_alloc(NENV * sizeof(*fxbufs));
2  memset(fxbufs, 0, NENV * sizeof(*fxbufs));
```

Then update the code of `trap` to save and store registers, in `trap.c`

```
1  if ((tf->tf_cs & 3) == 3) {
2      // Trapped from user mode.
3
4      assert(curenv);
5      lock_kernel();
6
7      // save registers
8      asm volatile ("fxsave %0::"m"(fxbufs[curenv->env_id]));
9
10     // .....
11 }
```

In `env.c`

```

1  assert(e->env_tf.tf_eflags & FL_IF);
2
3  asm volatile ("fxrstor %0:::m"(fxbufs[curenv->env_id]));
4
5  unlock_kernel();
6  env_pop_tf(&(curenv->env_tf));

```

I implemented a `yieldf.c` to test the fpu:

```

1  #include <inc/lib.h>
2
3  void
4  umain(int argc, char **argv) {
5      for(int i = 0; i < 3; i++) {
6          if(fork() == 0) break;
7      }
8      printf("Hello, I am environment %08x.\n", thisenv->env_id);
9      for (int i = 0; i < 3; i++) {
10         sys_yield();
11         printf("Back in environment %08x, iteration %f.\n",
12             thisenv->env_id, i * 1000.123123123);
13     }
14     printf("All done in environment %08x.\n", thisenv->env_id);
15 }

```

It works perfectly.

```

jos: zsh — Konsole
文件(F) 编辑(E) 视图(V) 书签(B) 设置(S) 帮助(H)
[00001000] new env 00001002
[00001000] new env 00001003
Hello, I am environment 00001000.
Hello, I am environment 00001001.
Hello, I am environment 00001002.
Hello, I am environment 00001003.
Back in environment 00001001, iteration 0.0.
Back in environment 00001002, iteration 0.0.
Back in environment 00001003, iteration 0.0.
Back in environment 00001000, iteration 0.0.
Back in environment 00001001, iteration 1000.123107910325633024.
Back in environment 00001002, iteration 1000.123107910325633024.
Back in environment 00001003, iteration 1000.123107910325633024.
Back in environment 00001000, iteration 1000.123123123000027590.
Back in environment 00001001, iteration 2000.246215820651266048.
All done in environment 00001001.
[00001001] exiting gracefully
[00001001] free env 00001001
Back in environment 00001000, iteration 2000.246246246000055180.
Back in environment 00001002, iteration 2000.246215820651266048.
All done in environment 00001002.
[00001002] exiting gracefully
[00001002] free env 00001002
All done in environment 00001000.
[00001000] exiting gracefully
[00001000] free env 00001000
Back in environment 00001003, iteration 2000.246215820651266048.
All done in environment 00001003.
[00001003] exiting gracefully
[00001003] free env 00001003
No runnable environments in the system!
Welcome to the JOS kernel monitor!
Type 'help' for a list of commands.
K> QEMU: Terminated
(base) → jos git:(lab4)

```

## Exercise 7

```

1  static envid_t
2  sys_exofork(void) {
3      struct Env * e;
4      ckret(env_alloc(&e, curenv->env_id));
5      e->env_status = ENV_NOT_RUNNABLE;
6      e->env_tf = curenv->env_tf;
7      e->env_tf.tf_regs.reg_eax = 0; // child returns 0
8      return e->env_id;
9  }

```

```

10
11 static int
12 sys_env_set_status(envid_t envid, int status) {
13     if(status < 0 || status >= ENV_STATUS_MAX)
14         return -E_INVALID;
15     int errno = 0;
16     struct Env *e;
17     if((errno = envid2env(envid, &e, true)) < 0) {
18         return errno;
19     }
20     e->env_status = status;
21     return 0;
22 }
23
24 static int
25 usr_mem_va_check(void *va) {
26     if((uintptr_t)va >= UTOP) return -E_INVALID;
27     if(PGOFf(va) != 0) return -E_INVALID;
28     return 0;
29 }
30
31 static int
32 usr_mem_perm_check(int perm) {
33     if((perm & PTE_SYSCALL) != perm) return -E_INVALID;
34     if(!(perm & PTE_U)) return -E_INVALID;
35     if(!(perm & PTE_P)) return -E_INVALID;
36     return 0;
37 }
38
39 static int
40 sys_page_alloc(envid_t envid, void *va, int perm) {
41     struct Env *e;
42     ckret(envid2env(envid, &e, true));
43     struct PageInfo *pg = page_alloc(0);
44     if(pg == NULL) return -E_NO_MEM;
45     ckret(usr_mem_perm_check(perm));
46     ckret(page_insert(e->env_pgdir, pg, va, perm));
47     return 0;
48 }
49
50 static int
51 sys_page_map(envid_t srcenvid, void *srcva, envid_t dstenvid, void *dstva, int perm){
52     struct Env *src, *dst;
53     ckret(envid2env(srcenvid, &src, true));
54     ckret(envid2env(dstenvid, &dst, true));
55     ckret(usr_mem_va_check(srcva));
56     ckret(usr_mem_va_check(dstva));
57     ckret(usr_mem_perm_check(perm));
58     pte_t *srcpte = pgdir_walk(src->env_pgdir, srcva, false);
59     if(srcpte == NULL) return -E_INVALID;
60     if((perm & PTE_W) && !(*srcpte & PTE_W)) return -E_INVALID;
61     ckret(page_insert(dst->env_pgdir, pa2page(PTE_ADDR(*srcpte)), dstva, perm));
62     return 0;
63 }
64
65 static int
66 sys_page_unmap(envid_t envid, void *va) {
67     struct Env *e;
68     ckret(envid2env(envid, &e, true));
69     page_remove(e->env_pgdir, va);
70     return 0;
71 }
72
73 int32_t
74 syscall(uint32_t syscallno, uint32_t a1, uint32_t a2, uint32_t a3, uint32_t a4, uint32_t a5) {
75     int32_t ret = 0;
76     switch (syscallno) {
77         case SYS_cgetc: ret = sys_cgetc(); break;
78         case SYS_cputs: sys_cputs((const char *)a1, a2); break;
79         case SYS_env_destroy: ret = sys_env_destroy(a1); break;
80         case SYS_getenvid: ret = sys_getenvid(); break;
81         case SYS_yield: sched_yield(); break;
82         case SYS_exofork: ret = sys_exofork(); break;
83         case SYS_env_set_status: ret = sys_env_set_status((envid_t)a1, (int)a2); break;
84         case SYS_page_alloc: ret = sys_page_alloc((envid_t)a1, (void*)a2, (int)a3); break;
85         case SYS_page_map: ret = sys_page_map((envid_t)a1, (void*)a2, (envid_t)a3, (void*)a4, (int)a5); break;
86         case SYS_page_unmap: ret = sys_page_unmap((envid_t)a1, (void*)a2); break;
87         case SYS_env_set_pgfault_upcall: ret = sys_env_set_pgfault_upcall((envid_t)a1, (void*)a2); break;
88         case SYS_ipc_try_send: ret = sys_ipc_try_send((envid_t)a1, (uint32_t)a2, (void*)a3, (unsigned int)a4); break;
89         case SYS_ipc_recv: ret = sys_ipc_recv((void*)a1); break;
90         default: ret = -E_INVALID; break;
91     }
92     return ret;
93 }

```

## Exercise 8

```

1 static int
2 sys_env_set_pgfault_upcall(envid_t envid, void *func) {
3     struct Env *e;
4     ckret(envid2env(envid, &e, true));
5     e->env_pgfault_upcall = func;
6     return 0;

```



```
7 }
```

## Exercise 9

```
1  if(curenv->env_pgfault_upcall) {
2      void * uesp = NULL;
3      if(UXSTACKTOP - PGSIZE <= tf->tf_esp && tf->tf_esp < UXSTACKTOP) {
4          uesp = (void*)(tf->tf_esp - 4);
5      }
6      else {
7          uesp = (void*)UXSTACKTOP;
8      }
9      struct UTrapframe * uxframe = (struct UTrapframe *) (uesp - sizeof(struct UTrapframe));
10     user_mem_assert(curenv, uxframe, sizeof(struct UTrapframe), PTE_P | PTE_U | PTE_W);
11     uxframe->utf_fault_va = fault_va;
12     uxframe->utf_err      = tf->tf_err;
13     uxframe->utf_regs     = tf->tf_regs;
14     uxframe->utf_eip      = tf->tf_eip;
15     uxframe->utf_eflags   = tf->tf_eflags;
16     uxframe->utf_esp      = tf->tf_esp;
17     curenv->env_tf.tf_esp = (uintptr_t)(uxframe);
18     curenv->env_tf.tf_eip = (uintptr_t)(curenv->env_pgfault_upcall);
19     env_run(curenv);
20 }
```

If the user environment runs out the exception stack, it will be destroyed.

## Exercise 10

```
1  _pgfault_upcall:
2      // Call the C page fault handler.
3      pushl %esp          // function argument: pointer to UTF
4      movl _pgfault_handler, %eax
5      call *%eax
6      addl $4, %esp       // pop function argument
7
8      // trap-time esp      48
9      // trap-time eflags   44
10     // trap-time eip      40
11     // utf_regs.reg_eax    36
12     // utf_regs.reg_ecx    32
13     // utf_regs.reg_edx    28
14     // utf_regs.reg_ebx    24
15     // utf_regs.reg_oesp   20
16     // utf_regs.reg_ebp    16
17     // utf_regs.reg_esi    12
18     // utf_regs.reg_edi     8
19     // utf_err (error code) 4
20     // utf_fault_va        <-- %esp 0
21
22     movl 40(%esp), %eax
23     movl 48(%esp), %ebx
24     subl $4, %ebx
25     movl %ebx, 48(%esp)
26     movl %eax, (%ebx)
27
28     // Restore the trap-time registers. After you do this, you
29     // can no longer modify any general-purpose registers.
30     // LAB 4: Your code here.
31
32     addl $8, %esp
33     popal
34     // trap-time esp      8
35     // trap-time eflags   4
36     // trap-time eip      <-- %esp 0
37
38     // Restore eflags from the stack. After you do this, you can
39     // no longer use arithmetic operations or anything else that
40     // modifies eflags.
41     // LAB 4: Your code here.
42
43     addl $4, %esp
44     popfl
45
46     // Switch back to the adjusted trap-time stack.
47     // LAB 4: Your code here.
48
49     popl %esp
50
51     // Return to re-execute the instruction that faulted.
52     // LAB 4: Your code here.
53
54     ret
```

## Exercise 11

```

1 void
2 set_pgfault_handler(void (*handler)(struct UTrapframe *utf))
3 {
4     int r;
5
6     if (_pgfault_handler == 0) {
7         // First time through!
8         // LAB 4: Your code here.
9         r = sys_page_alloc(0, (void*)(UXSTACKTOP - PGSIZE), PTE_P | PTE_U | PTE_W);
10        if(r < 0) panic("sys_page_alloc fail %e", r);
11        r = sys_env_set_pgfault_upcall(0, _pgfault_upcall);
12        if(r < 0) panic("set_pgfault_handler fail %e", r);
13        // panic("set_pgfault_handler not implemented");
14    }
15
16    // Save handler pointer for assembly to call.
17    _pgfault_handler = handler;
18 }

```

## Exercise 12

```

1 static void
2 pgfault(struct UTrapframe *utf)
3 {
4     void *addr = (void *) utf->utf_fault_va;
5     uint32_t err = utf->utf_err;
6     int r;
7
8     assert_panic(err & FEC_PR, "page fault: access non-present page: %p, eip=%08x", addr, utf->utf_eip);
9     assert_panic(err & FEC_U, "page fault: no permission to access: %p, eip=%08x", addr, utf->utf_eip);
10    assert_panic(err & FEC_W, "page fault: not readable: %p, eip=%08x", addr, utf->utf_eip);
11
12    pte_t pte = get_pte(PDX(addr), PTX(addr));
13    assert_panic(is_masked(pte, PTE_P | PTE_U | PTE_COW), "page fault: not writable %p, eip=%08x", addr, utf->utf_eip);
14
15    void * page_addr = ROUNDDOWN(addr, PGSIZE);
16    assert(sys_page_alloc(0, UTEMP, PTE_P | PTE_U | PTE_W) == 0);
17    memmove(UTEMP, page_addr, PGSIZE);
18    assert(sys_page_map(0, UTEMP, 0, page_addr, PTE_P | PTE_U | PTE_W) == 0);
19    assert(sys_page_unmap(0, UTEMP) == 0);
20 }
21
22 duppage(envid_t envid, void * pageaddr)
23 {
24     // LAB 4: Your code here.
25    ckret(sys_page_map(0, pageaddr, envid, pageaddr, PTE_P | PTE_U | PTE_COW));
26    ckret(sys_page_map(0, pageaddr, 0, pageaddr, PTE_P | PTE_U | PTE_COW));
27    return 0;
28 }
29
30 envid_t
31 fork(void)
32 {
33     // LAB 4: Your code here.
34     // panic("fork not implemented");
35    set_pgfault_handler(pgfault);
36
37    envid_t envid = sys_exofork();
38    if(!envid) {
39        thisenv = envs + ENVX(sys_getenvid()); // this is very important
40        return envid;
41    }
42
43    for(size_t i = 0; i < PDX(UTOP); i++) {
44        pde_t pde = get_pde(i);
45        if(!is_masked(pde, PTE_P)) continue;
46        for(size_t j = 0; j < NPTENTRIES; j++) {
47            pte_t pte = get_pte(i, j);
48            void * pgaddr = PGADDR(i, j, 0);
49            if(pgaddr == (void*)(UXSTACKTOP - PGSIZE)) // ignore UXSTACK
50                continue;
51            if(!is_masked(pte, PTE_P | PTE_U)) // user not accessible
52                continue;
53            if(is_masked(pte, PTE_W) || is_masked(pte, PTE_COW)) { // COW page
54                ckret(duppage(envid, pgaddr));
55            }
56            else{
57                ckret(sys_page_map(0, pgaddr, envid, pgaddr, PTE_FLAGS(pte)));
58            }
59        }
60    }
61
62    ckret(sys_page_alloc(envid, (void*)(UXSTACKTOP - PGSIZE), PTE_P | PTE_U | PTE_W));
63
64    extern void _pgfault_upcall(void);
65    ckret(sys_env_set_pgfault_upcall(envid, _pgfault_upcall));
66
67    ckret(sys_env_set_status(envid, ENV_RUNNABLE));
68    return envid;
69 }

```

## Exercise 13

A new entry is added into the `kern/trapentry.s`. (The initialization in `kern/trap.c` is automatically setup).

```
1 TRAPHANDLER_NOEC( trap_handler_SYSCALL , T_SYSCALL , PL_USER )
```

## Exercise 14

```
1 switch (tf->tf_trapno) {
2     case T_DEBUG: monitor(tf); break;
3     case T_PGFLT: page_fault_handler(tf); break;
4     case T_BRKPT: monitor(tf); break;
5     case T_SYSCALL:
6         // The system call number will go in %eax,
7         // and the arguments (up to five of them) will go in %edx, %ecx, %ebx, %edi, and %esi, respectively.
8         // The kernel passes the return value back in %eax.
9         tf->tf_regs.reg_eax = syscall(
10             tf->tf_regs.reg_eax,
11             tf->tf_regs.reg_edx, tf->tf_regs.reg_ecx, tf->tf_regs.reg_ebx,
12             tf->tf_regs.reg_edi, tf->tf_regs.reg_esx
13         );
14         env_run(curenv);
15         break;
16     case IRQ_OFFSET + IRQ_TIMER:
17         // tell the lapic the interrupt has finished
18         // for timers, no matter where the eoi signal is,
19         // the next interrupt won't be influenced.
20         lapic_eoi();
21         sched_yield();
22         break;
23     default: break;
24 }
```

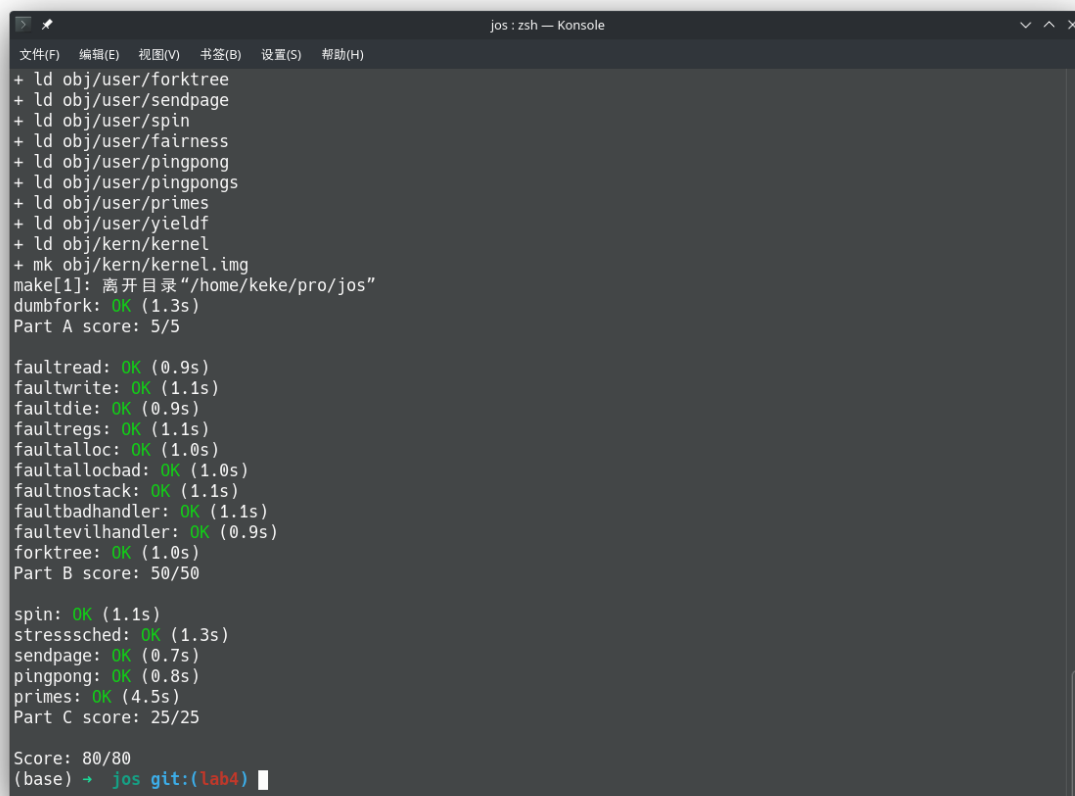
## Exercise 15

```
1 static int
2 sys_ipc_try_send(envid_t envid, uint32_t value, void *srcva, unsigned perm) {
3     struct Env *dst;
4     ckret(envid2env(envid, &dst, false));
5     astret(dst->env_ipc_recving, -E_IPC_NOT_RECV);
6     dst->env_ipc_recving = 0;
7     dst->env_ipc_from = curenv->env_id;
8     if((uintptr_t)srcva < UTOP && (uintptr_t)dst->env_ipc_dstva < UTOP) {
9         ckret(usr_mem_va_check(srcva));
10        pte_t *pte = pgdir_walk(curenv->env_pgdir, srcva, false);
11        astret(pte, -E_INVALID);
12        astret(is_masked(*pte, PTE_P), -E_INVALID);
13        ckret(usr_mem_perm_check(perm));
14        if((perm & PTE_W) && !(*pte & PTE_W)) return -E_INVALID;
15        ckret(page_insert(dst->env_pgdir, pa2page(PTE_ADDR(*pte)), dst->env_ipc_dstva, perm));
16        dst->env_ipc_perm = perm;
17    }
18    else {
19        dst->env_ipc_perm = 0;
20    }
21    dst->env_ipc_value = value;
22    dst->env_status = ENV_RUNNABLE;
23    return 0;
24 }
25
26 static int
27 sys_ipc_recv(void *dstva) {
28     struct Env *e = curenv;
29     if((uintptr_t)dstva < UTOP) {
30         ckret(usr_mem_va_check(dstva));
31     }
32     e->env_ipc_dstva = dstva;
33     e->env_ipc_value = 0;
34     e->env_tf.tf_regs.reg_eax = 0;
35     e->env_ipc_recving = 1;
36     e->env_status = ENV_NOT_RUNNABLE;
37     sched_yield();
38     return 0;
39 }
```

```
1 int32_t
2 ipc_recv(envid_t *from_env_store, void *pg, int *perm_store)
3 {
```

```
4 // LAB 4: Your code here.
5 // panic("ipc_rcv not implemented");
6 int r = sys_ipc_rcv(pg ? pg : (void*)UTOP);
7 if(r < 0) {
8     if(from_env_store) *from_env_store = 0;
9     if(perm_store) *perm_store = 0;
10    // logd("rcv error");
11    return r;
12 }
13 else {
14     if(from_env_store) *from_env_store = thisenv->env_ipc_from;
15     if(perm_store) *perm_store = thisenv->env_ipc_perm;
16     // logd("rcv: %d", thisenv->env_ipc_value);
17     return thisenv->env_ipc_value;
18 }
19 }
20
21 void
22 ipc_send(envid_t to_env, uint32_t val, void *pg, int perm)
23 {
24     // LAB 4: Your code here.
25     // panic("ipc_send not implemented");
26     pg = pg ? pg : (void*) UTOP;
27     while(sys_ipc_try_send(to_env, val, pg, perm) < 0) {
28         asm volatile("pause");
29         sys_yield();
30     }
31 }
```

## The end of this lab



```
jos: zsh — Konsole
文件(F) 编辑(E) 视图(V) 书签(B) 设置(S) 帮助(H)
+ ld obj/user/forktree
+ ld obj/user/sendpage
+ ld obj/user/spin
+ ld obj/user/fairness
+ ld obj/user/pingpong
+ ld obj/user/pingpongs
+ ld obj/user/primes
+ ld obj/user/yieldf
+ ld obj/kern/kernel
+ mk obj/kern/kernel.img
make[1]: 离开目录"/home/keke/pro/jos"
dumbfork: OK (1.3s)
Part A score: 5/5

faultread: OK (0.9s)
faultwrite: OK (1.1s)
faultdie: OK (0.9s)
faultregs: OK (1.1s)
faultalloc: OK (1.0s)
faultallocbad: OK (1.0s)
faultnostack: OK (1.1s)
faultbadhandler: OK (1.1s)
faultevilhandler: OK (0.9s)
forktree: OK (1.0s)
Part B score: 50/50

spin: OK (1.1s)
stresssched: OK (1.3s)
sendpage: OK (0.7s)
pingpong: OK (0.8s)
primes: OK (4.5s)
Part C score: 25/25

Score: 80/80
(base) -> jos git:(lab4)
```